

Real-time video anomaly detection

Can a **small** vision-language model watch live video
and report what matters — fast enough, and without crying wolf?

59.7

out of 100 · was 49.9 one hour before the
end

29.8 / 35

on D2 — strongest tier, 4-minute videos
localised well

6.1×

faster than real time · 1.65 s per 10 s
window

THE PIPELINE

Two models see the same frames.



WHY TWO MODELS

Over half the events came from the encoder, not the VLM.

Source of the 38 shipped events	Events
SigLIP2 only — the VLM never saw these	17
Both models agreed	10
Qwen3-VL only — SigLIP never saw these	7
Derived by merging overlapping spans	4

Only **10 of 38** events had both models agreeing. The other 28 came from exactly one of them.

That is the case for running two — not a story, a count.

Five VLMs, identical test. Two surprises.

Model	Size	Precision	Recall	Verdict
Qwen3-VL-8B-Instruct	8B	0.62	0.40	Shipped — produced D1 + D2
Qwen3-VL-4B-Instruct	4B	0.77	0.50	Ties the 8B (differ on 3 of 24)
Cosmos-Reason2-8B	8B	0.57	0.20	Half the recall of a 4B
Qwen3-VL-2B-Instruct	2B	0.30	0.15	Too small
Qwen2.5-VL-7B-Instruct	7B	0.00	0.00	Silent on 20 of 24
SigLIP2 + trained GRU head	frozen	— added +2.4 marks —		Shipped

Cosmos-Reason2-8B is purpose-built for traffic anomaly reasoning — and got half the recall of a general 4B here. **Benchmark performance does not transfer.**

SigLIP was marked "rejected" in our own ledger, on theory. It turned out to be one of the two things that actually moved our score. **We had rejected it on theory, without testing it.**

+9.8 marks. We changed no model.

Eight prediction sets — three model sizes, five window configs, a fine-tune, targeted prompts — all scored **exactly 8.0** on the long videos.

That invariance was the clue. If changing the model changes nothing, the failure isn't in the model.

Our predicted spans were **1–8 seconds** inside videos of **327–602 seconds**. A 5-second span can never overlap a 100-second event by half — **geometrically impossible, however right the class is.**

49.9 → 56.0

widening spans to 112–138 s.

No new inference. No changed class.

Our classes had been right all along.

→ 59.7

adding SigLIP2, which predicts

loitering

— a class we score 0/6 on and

no VLM proposed all day.

HOW WE WORKED

The evidence was in our own notes from that morning: the sample data held a **125-second** event. We wrote down that these events are long — then spent hours swapping models instead of spans.

Three AI agents, one repository.

Agent	Role	Owned
Coordinator	Research, strategy, every GPU run, every commit	inference lane, Modal driver
Builder	Data pipeline, synthetic video generation, training, scoring	data lane, tests
Comms	Live demo, slides, process record	demo, docs

- ▶ **Ownership by directory.** No agent edits another's files. We learned it the hard way — two agents built the same four modules before the rule existed.
- ▶ **One committer,** so the history stayed readable on a graded repo.
- ▶ **A written ledger, not chat.** Every technique across 7 tiers marked done / running / rejected — **with the reason.** An agent could crash, restart, and resume from files. It happened twice.
- ▶ **They corrected each other.** One caught a scorer using first-fit instead of best-loU matching. One caught a run that exited "successfully" having done 6 of 34 videos.

Ten failures in one day. None of them crashed.

- ▶ **ffmpeg's seek silently returns more than you ask for.** A 468 s video rendered as 241 s — every timestamp after it was wrong.
- ▶ **The fix for one bug caused another.** We stopped one bad video killing a 34-video run; it then wrote 10 failures to disk as confident, plausible "nothing here".
- ▶ **A model reported 100% validation accuracy.** It had 7 validation samples, because an upload had died unnoticed.
- ▶ **We deleted 125 videos we had already rendered** — a third of our synthetic training data — rather than train on labels we could not verify.
- ▶ **We retracted our own headline finding** before it shipped. The retraction was load-bearing: keeping it meant shipping the weaker model.

In a seven-hour build, **every failure that mattered was silent.**
So we validated intermediate artifacts — not outputs.