

Demo brief — run sheet

`http://127.0.0.1:8800` · an operations console: video left, alerts right, scoreboard bottom, and a button that **runs the real model on an A100 while the audience watches**.

Before you present — 3 minutes

```
.venv/bin/modal deploy demo/modal_live.py    # prints the URL
# put it in .env as MODAL_INFER_URL=...
.venv/bin/python demo/server.py             # http://127.0.0.1:8800
```

1. Open `/health` — confirms the backends are live.
 2. **Press Run once, privately.** First call pays a **~50 s cold start** while the model loads. After that a window takes **~9 s**. Never let the audience watch the cold start.
 3. Afterwards: `.venv/bin/modal app stop ahc-demo-live`.
-

The 4 minutes on stage

#	Do	Say
1	Play a short clip with a clean hit	"Frames sampled by timestamp, 20-second window, one call."
2	Point at the timeline	"Top row is what the model was shown. Middle is truth. Bottom is what it reported."
3	Press Enter on a window	"That's not a recording — that's an A100 answering right now."
4	Point at the timing breakdown	"1.8 seconds of model time against a 10-second window. 5.5x real time."
5	Show a false alarm (red)	"We show our misses. False alarms are the thing this problem punishes."
6	Scoreboard	"59.7 out of 100. D2 is our strongest tier."

Keyboard: `space` play/pause · `↔` 5 s · `Enter` runs the model.

The honesty features — worth pointing out

- **Every alert says live or recorded.** A pill in the header names which backend answered. Nothing on the page implies GPU work that did not happen.
- **The live endpoint reuses `vad.frames.sample` and `vad.prompts` verbatim** — the same 16 frames over the same 20-second window as the scored runs. It is the same system, not a demo mock-up.
- **Three states are never blurred:** no prediction = *not run* · `"failed"` = the run broke · `"events": []` = the model genuinely said normal, and *is* scored.
- Gemini is available as a fallback but labelled **amber, not runtime-legal** — a hosted model cannot be in the submission's runtime path.

If it breaks

- **Button hangs** → per-backend deadlines fire (Modal 120 s, Gemini 12 s). It fails and says why.
- **Modal down** → falls back to the recorded run and states it on screen. The demo still works.
- **Video won't play** → you must use `demo/server.py`, not `python -m http.server`, which ignores HTTP Range and hangs the player.

Other pages

`/review` — all 34 videos, every prediction scored · `/slides` — the two submission slides