

Two models see the same frames.



WHY TWO MODELS

Over half the events came from the encoder, not the VLM.

submission JSON → 59.7 / 100

Source of the 38 shipped events	Events
SigLIP2 only — the VLM never saw these	17
Both models agreed	10
Qwen3-VL only — SigLIP never saw these	7
Derived by merging overlapping spans	4

Only **10 of 38** events had both models agreeing. The other 28 came from exactly one of them.
That is the case for running two — not a story, a count.

Union, and an asymmetric span policy.

1 · Union, never replace

- ▶ If either model names a class the other missed, it goes in.
- ▶ SigLIP found **loitering** on two long videos — a class we score **0 / 6** on, and the VLM never proposed it once, all day.
- ▶ Replacing one model with the other scored **lower**. Both times.

Latency, measured

- ▶ **1.65 s** mean per window · p50 **0.96 s** · p95 **3.20 s**
- ▶ A window closes every 10 s → **6.1× real time**, 3.1× at p95, on one A100.

2 · Span policy is per-difficulty

- ▶ **4-minute videos**: spans left as the model reported them, 6–30 s.
- ▶ **6–10 minute videos**: same-class events merged and padded to **232–602 s**.
- ▶ We tested widening the 4-minute spans the same way. **It scored worse**.
- ▶ So short events really are short and long events really are long — **a measured property of the data, not a tuning trick**.

Widening the long-video spans was worth **+6.1 marks** on its own — with **no new inference and no changed class**. The classes had been right all along; the geometry was wrong.