


The training data is short clips. The marks are in long video.

Real-time video anomaly detection · 11 classes
Qwen3-VL-4B-Instruct · one A100-40GB
34 public test videos · 56 minutes of footage

Every training clip is 5–30 s with one event. We are scored on 240–629 s with up to six.


Drawn to a single time scale. Green blocks are the events a model must find and localise to $\pm$ IoU 0.5.



75 of the 100 marks had no training data at all — so we manufactured it: 360 long multi-event videos with exact ground truth.

Pipeline

Everything cacheable is cached, so tuning the output policy costs no GPU.



10 s

WINDOW CLOSES EVERY

1.83 s

TO SCORE IT

5.5×

FASTER THAN REAL TIME


\$0

TO TUNE THE POLICY

2.8× at p95 — real time holds on the tail. 326 windows, 5,216 frames, one A100-40GB.

Why a 4B

Five VLMs, identical 24-video set, same prompt and policy. F1.




The 4B and 8B are **statistically indistinguishable** — they differ on 3 videos of 24, $p = 0.25$. We ship the 4B because it is **half the size at equal measured accuracy**.

Nothing crashed. That was the problem.

All numbers measured on the 34-video public set
24 videos, 20 positives → ±1 TP ≈ 5 pp of recall
applied to every number on this page

Where the marks actually are


First scored upload. D2's 11.7 is exactly the empty-submission floor.



10× One D3 video is worth **ten** D1 videos — 10.00 vs 1.04 marks each — and we are at zero on D3.

Lower false-alarm, for free

One entrant on this leaderboard lost 26 of 35 marks to 94 false alarms.



AND THE OBVIOUS FIX THAT DOES NOT WORK

Gating on the model's own confidence, 0.5 → 0.9, removes **2 false alarms out of 31**. Confidence carries no signal — whatever gates false alarms must come from outside the model.

Seven silent failures in seven hours

None crashed. Every one would have read as “the model just isn't very good.”

- 1 `ffmpeg -t` silently returns a shorter segment
- 2 49 of 2,092 train rows have `end ≤ start`
- 3 clip pool keyed on folder, not class label
- 4 empty output ≡ a correct “no events”
- 5 one bad API reply cut a 34-video run to 6, exit 0
- 6 buffered output hid a 20-minute GPU job
- 7 the fix for #5 caused #7 – failures written as `[]`

So we validate intermediate artifacts, not outputs. #7 was caught only by the diagnostic added for #4. Each is now a test.

WE RETRACTED OUR OWN HEADLINE, BEFORE IT SHIPPED

“Bigger is worse — the 4B beats the 8B, F1 0.61 vs 0.48.” The two models differ on **3 videos out of 24**; sign test **p = 0.25**. The gap was 10 TP vs 8 TP — inside the ±1 TP noise band we had been telling everyone to apply, and then failed to apply ourselves.

A NEGATIVE RESULT, TESTED PROPERLY

0.556 AUC · CHANCE IS 0.500

Codec motion vectors as an ego-motion-robust gate — free from the decoder, the idea every fixed-camera paper implies. Against a 2,000-trial permutation null it **does not work**. Settled in 1,165 CPU-seconds, no GPU.

PLACEHOLDER · FINE-TUNE

The LoRA run had not finished at build time, so **no fine-tuned number appears anywhere on these slides**. If it lands, it goes in the model table on slide 1 and the marks chart here — and nowhere else.